

Artificial Intelligence II (CS4442B & CS9542B)

Supervised Learning: Linear Regression

Boyu Wang
Department of Computer Science
University of Western Ontario

Motivation examples

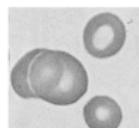
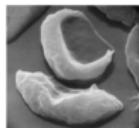
Classification



X = Document



Sports
Science
News



Y = Topic



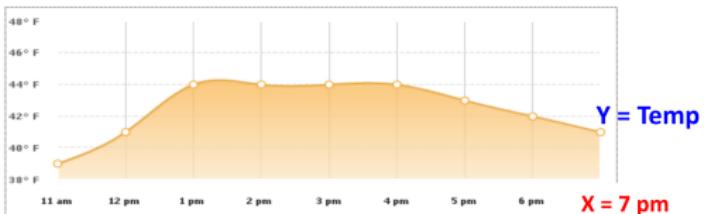
Anemic cell
Healthy cell

X = Cell Image

Y = Diagnosis

Motivation examples

Weather Prediction

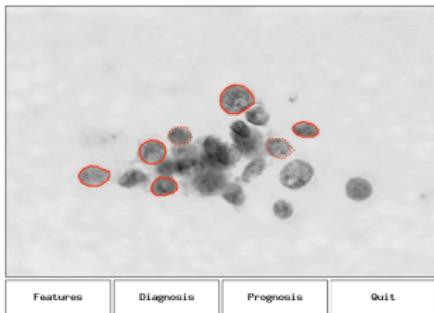


Estimating Contamination



Slide credit: Maria-Florina Balcan

Motivation examples



- Cell samples were taken from tumors in breast cancer patients before surgery, and imaged
- Tumors were excised
- Patients were followed to determine whether or not the cancer recurred, and how long until recurrence or disease free

Slide credit: Doina Precup

Data

- Thirty real-valued variables per tumor.
- Two variables that can be predicted:
 - Outcome (R=recurrence, N=non-recurrence)
 - Time (until recurrence, for R, time healthy, for N).

tumor size	texture	perimeter	...	outcome	time
18.02	27.6	117.5		N	31
17.99	10.38	122.8		N	61
20.29	14.34	135.1		R	27
...					

Slide credit: Doina Precup

Terminology

tumor size	texture	perimeter	...	outcome	time
18.02	27.6	117.5		N	31
17.99	10.38	122.8		N	61
20.29	14.34	135.1		R	27
...					

- ▶ Columns are called **features** or attributes or input variables.
- ▶ The outcome and time (which we are trying to predict) are called **output** variables or targets.
- ▶ A row in the table is called **training example** or **instance**.
- ▶ The problem of predicting the recurrence is called (binary) **classification**.
- ▶ The problem of predicting the time is called **regression**.

Problem formulation

tumor size	texture	perimeter	...	outcome	time
18.02	27.6	117.5		N	31
17.99	10.38	122.8		N	61
20.29	14.34	135.1		R	27
...					

- ▶ A training example i has the form: (x_i, y_i) , where $x_i \in \mathbb{R}^n$ is the number of features (feature dimension), and $y_i \in \{0, 1\}$ for (binary) classification, and $y_i \in \mathbb{R}$ for regression.
- ▶ The training set consists of m examples: $S = \{(x_i, y_i)\}_{i=1}^m$.
- ▶ We denote the matrix of features by $X = [x_1, \dots, x_m]^\top \in \mathbb{R}^{m \times n}$, and the column vector of outputs by $Y \in \mathbb{R}^m$.
- ▶ **Objective:** construct a **good predictor** using training data S (i.e., X and Y), such that it works well for **new patients (test data)**.

Problem formulation

- ▶ Let \mathcal{X} denote the space of input values
- ▶ Let \mathcal{Y} denote the space of output values
- ▶ Given a data set $S \in \mathcal{X} \times \mathcal{Y}$, find a function:

$$h : \mathcal{X} \rightarrow \mathcal{Y}$$

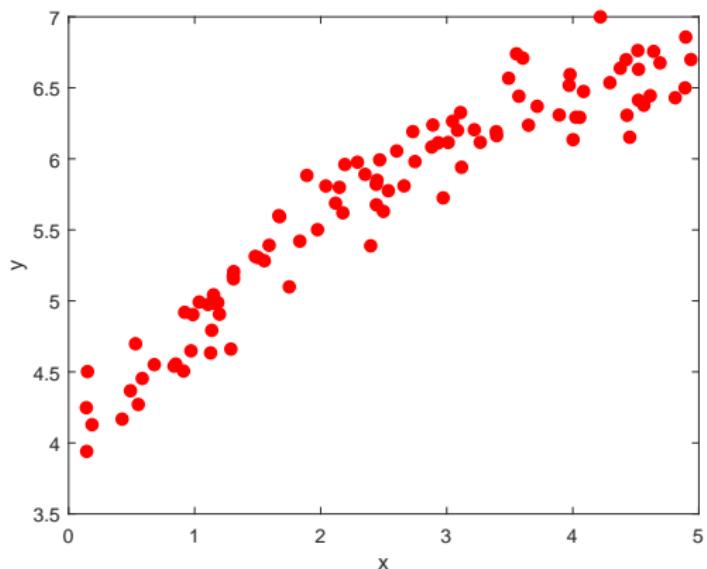
such that for a new example (x, y) , $h(x)$ can correctly predict the value of y .

- ▶ Types of supervised learning problem: classification, regression, ranking, structured prediction...
- ▶ **Key assumption:** training and test data are sampled from the same distribution.

Steps to solving a supervised learning problem

- ▶ Collect the data set (input-output pairs)
- ▶ Choose a class of hypotheses (hypothesis space) \mathcal{H}
- ▶ Choose a hypothesis $h \in \mathcal{H}$

Example: what hypothesis class should we choose?



Steps to solving a supervised learning problem

- ▶ Collect the data set (input-output pairs)
- ▶ Choose a class of hypotheses (hypothesis space) \mathcal{H}
- ▶ Choose a hypothesis $h \in \mathcal{H}$

Linear model

- ▶ In linear regression, we consider the model h_w has the form:

$$h_w(x) = w^\top x + b$$

where $w = [w_1, \dots, w_n]^\top \in \mathbb{R}^n$, w_i 's are called **parameters** or **weights**, and b is called **bias** or **intercept** term.

- ▶ To simplify the notation, we can add a feature to the other n features of x : $x \rightarrow [1; x] \in \mathbb{R}^{n+1}$, such that we will have

$$h_w(x) = w^\top x$$

where w and x are vectors of size $n + 1$.

- ▶ How to choose w ?

Error (cost) Minimization

- ▶ Intuitively, w should make the predictions of h_w close to the true values y on the data we have.
- ▶ Therefore, we will define an **error function** or cost function to measure the difference between our prediction and the true value
- ▶ We will choose w such that the error function is minimized:

$$w = \arg \min_w \sum_{i=1}^m \ell(h_w(x_i), y_i), \quad (1)$$

where ℓ is a loss function.

- ▶ (1) is the fundamental problem in machine learning!
 - How to choose ℓ
 - How to solve (1)

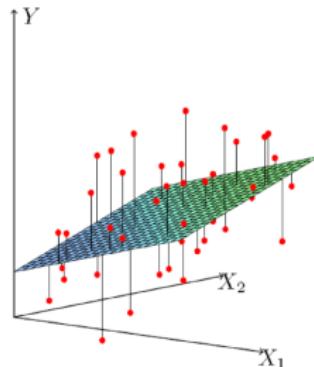
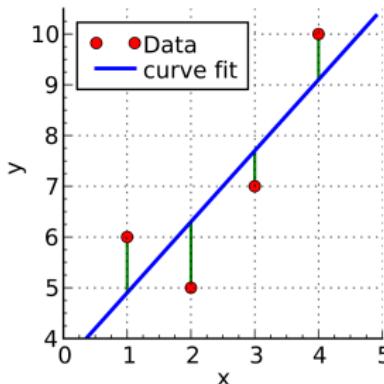
Least squares

- We use the squared error to measure our prediction performance, and choose w by minimizing the sum-of-squared errors:

$$J(w) = \frac{1}{2} \sum_{i=1}^m (h_w(x_i) - y_i)^2$$

(the $\frac{1}{2}$ is just for convenience)

- Fit by solving $\min_w J(w)$



Solve $\nabla J(w) = 0$ using the definition

- In order to minimize $J(w)$, we need to solve $\nabla J(w) = 0$

- $\nabla J(w) = [\frac{\partial J(w)}{w_1}, \dots, \frac{\partial J(w)}{w_n}]^\top$

$$\begin{aligned}\frac{\partial J(w)}{w_j} &= \frac{\partial}{\partial w_j} \frac{1}{2} \sum_{i=1}^m (h_w(x_i) - y_i)^2 \\ &= \frac{1}{2} \cdot 2 \sum_{i=1}^m (h_w(x_i) - y_i) \frac{\partial}{\partial w_j} (h_w(x_i) - y_i) \\ &= \sum_{i=1}^m (h_w(x_i) - y_i) \frac{\partial}{\partial w_j} (w^\top x_i - y_i) \\ &= \sum_{i=1}^m (h_w(x_i) - y_i) x_{i,j}\end{aligned}$$

where $x_{i,j}$ is the j -th feature of the i -th instance.

- Setting all these partial derivatives to 0, we get a linear system with $(n+1)$ equations and $(n + 1)$ unknown variables.

Solve $\nabla J(w) = 0$ using vector calculus

$$\begin{aligned}\nabla J(w) &= \nabla_w \left(\frac{1}{2} \sum_{i=1}^m (h_w(x_i) - y_i)^2 \right) \\&= \nabla_w \left(\frac{1}{2} (Xw - y)^\top (Xw - y) \right) \\&= \nabla_w \frac{1}{2} \left(w^\top X^\top Xw - y^\top Xw - w^\top X^\top y + y^\top y \right) \\&= X^\top Xw - X^\top y\end{aligned}$$

- ▶ Setting $\nabla J(w) = 0$ gives

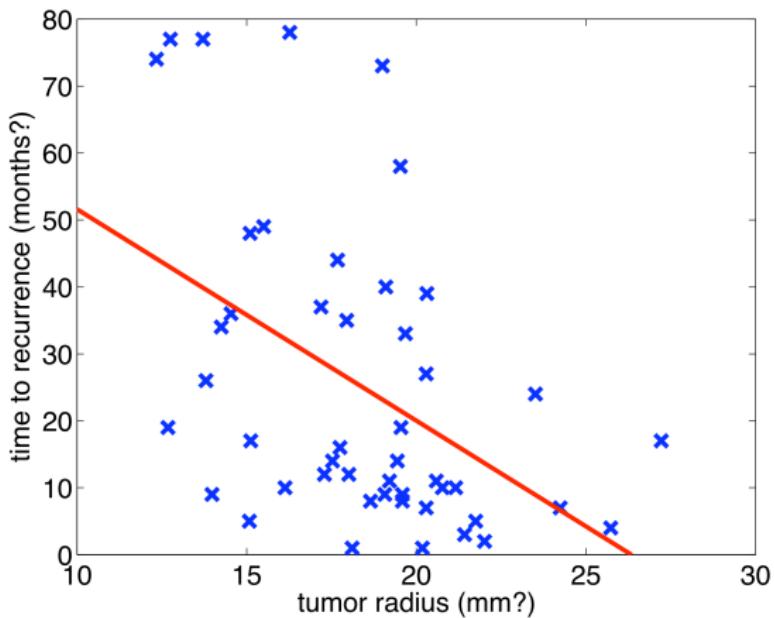
$$\begin{aligned}X^\top Xw - X^\top y &= 0 \\ \Rightarrow X^\top Xw &= X^\top y \\ \Rightarrow w &= (X^\top X)^{-1} X^\top y\end{aligned}$$

- ▶ The inverse exists if the columns of X are linearly independent.

Linear regression summary

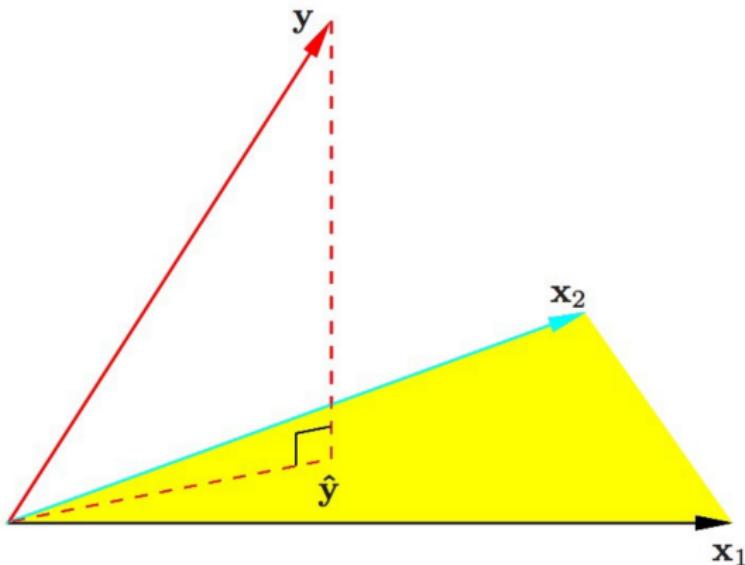
- ▶ The solution is $w = (X^T X)^{-1} X^T y$, where X is the data matrix augmented with a column of ones, and y is the column vector of target outputs.
- ▶ The number of data points should be at least larger or equal to the number of features to ensure a unique solution exists.
- ▶ An analytical, exact solution exists for linear regression, which is a rare case in machine learning.
- ▶ Demo!

Predicting recurrence time based on tumor size



Slide credit: Doina Precup

Geometric interpretation orthogonal projection



<https://towardsdatascience.com/3-things-you-didnt-know-about-simple-linear-regression-b5a86362dd53>

Probabilistic perspective of linear regression

- Assume y_i is a noisy target value, generated from a hypothesis $h_{\mathbf{w}}(\mathbf{x})$
- More specifically, assume that there exists \mathbf{w} such that:

$$y_i = h_{\mathbf{w}}(\mathbf{x}_i) + \epsilon_i$$

where ϵ_i is random variable (noise) drawn independently for each \mathbf{x}_i according to some Gaussian (normal) distribution with mean zero and variance σ .

- How should we choose the parameter vector \mathbf{w} ?

Bayes rule in machine learning

Let h be a hypothesis and D be the set of training data.
Using Bayes theorem, we have:

$$P(h|D) = \frac{P(D|h)P(h)}{P(D)},$$

where:

- $P(h)$ is the *prior probability of hypothesis h*
- $P(D) = \int_h P(D|h)P(h)$ is the probability of training data D (normalization, independent of h)
- $P(h|D)$ is the probability of h given D
- $P(D|h)$ is the probability of D given h (*likelihood of the data*)

Select a hypothesis by maximum a posteriori (MAP)

- What is the most probable hypothesis given the training data?
- *Maximum a posteriori (MAP)* hypothesis h_{MAP} :

$$\begin{aligned} h_{MAP} &= \arg \max_{h \in \mathcal{H}} P(h|D) \\ &= \arg \max_{h \in \mathcal{H}} \frac{P(D|h)P(h)}{P(D)} \text{(using Bayes theorem)} \\ &= \arg \max_{h \in \mathcal{H}} P(D|h)P(h) \end{aligned}$$

Last step is because $P(D)$ is independent of h (so constant for the maximization)

- This is the Bayesian answer (more in a minute)

Maximum likelihood estimation

$$h_{MAP} = \arg \max_{h \in \mathcal{H}} P(D|h)P(h)$$

- If we assume $P(h_i) = P(h_j)$ (all hypotheses are equally likely a priori) then we can further simplify, and choose the *maximum likelihood (ML) hypothesis*:

$$h_{ML} = \arg \max_{h \in \mathcal{H}} P(D|h) = \arg \max_{h \in \mathcal{H}} L(h)$$

- Standard assumption: the training examples are *independently identically distributed (i.i.d.)*
- This allows us to simplify $P(D|h)$:

$$P(D|h) = \prod_{i=1}^m P(\langle \mathbf{x}_i, y_i \rangle | h) = \prod_{i=1}^m P(y_i | \mathbf{x}_i; h)P(\mathbf{x}_i)$$

The log trick

- We want to maximize:

$$L(h) = \prod_{i=1}^m P(y_i|\mathbf{x}_i; h)P(\mathbf{x}_i)$$

This is a product, and products are hard to maximize!

- Instead, we will maximize $\log L(h)$! (the log-likelihood function)

$$\log L(h) = \sum_{i=1}^m \log P(y_i|\mathbf{x}_i; h) + \sum_{i=1}^m \log P(\mathbf{x}_i)$$

- The second sum depends on D , but not on h , so it can be ignored in the search for a good hypothesis

Maximum likelihood for regression

- Adopt the assumption that:

$$y_i = h_{\mathbf{w}}(\mathbf{x}_i) + \epsilon_i,$$

where $\epsilon_i \sim \mathcal{N}(0, \sigma)$.

- The best hypothesis maximizes the likelihood of $y_i - h_{\mathbf{w}}(\mathbf{x}_i) = \epsilon_i$
- Hence,

$$L(\mathbf{w}) = \prod_{i=1}^m \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2}\left(\frac{y_i - h_{\mathbf{w}}(\mathbf{x}_i)}{\sigma}\right)^2}$$

because the noise variables ϵ_i are from a Gaussian distribution

Applying the log trick

$$\begin{aligned}\log L(\mathbf{w}) &= \sum_{i=1}^m \log \left(\frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2} \frac{(y_i - h_{\mathbf{w}}(\mathbf{x}_i))^2}{\sigma^2}} \right) \\ &= \sum_{i=1}^m \log \left(\frac{1}{\sqrt{2\pi\sigma^2}} \right) - \sum_{i=1}^m \frac{1}{2} \frac{(y_i - h_{\mathbf{w}}(\mathbf{x}_i))^2}{\sigma^2}\end{aligned}$$

Maximizing the right hand side is the same as minimizing:

$$\sum_{i=1}^m \frac{1}{2} \frac{(y_i - h_{\mathbf{w}}(\mathbf{x}_i))^2}{\sigma^2}$$

This is our old friend, the sum-squared-error function! (the constants that are independent of h can again be ignored)

Slide credit: Doina Precup

Maximum likelihood and linear regression

- Under the assumption that the training examples are i.i.d. and that we have Gaussian target noise, the maximum likelihood solution is equivalent to the least-squares solution for linear regression:

$$w = \arg \min_w \sum_{i=1}^m (h_w(x_i) - y_i)^2$$

- If the noise is not normally distributed, maximizing the likelihood will not be the same as minimizing the sum-squared error
- In practice, different loss functions are used depending on the noise assumption.
- Squared loss is the most popular one because it is the simplest one!

Extending linear regression to more complex models

- ▶ Linear regression should be the first thing you try for real-valued outputs!
- ▶ Two possible solutions:
 1. Transform the data – apply a (nonlinear) transformation $x \rightarrow \phi(x)$, then do linear regression in the transformed space. (our focus in this lecture)
 - ▶ e.g., log, exp, sin, cos
 - ▶ polynomial transformation: cross-terms, higher-order terms...
 - ▶ basis expansions: Gaussian basis, polynomial basis...
 - ▶ Dummy coding of categorical features
 - ▶ Use a different hypothesis class (e.g. non-linear functions)

Linear models in general

- Given a set of examples $\langle \mathbf{x}_i, y_i \rangle_{i=1 \dots m}$, we fit a hypothesis

$$h_{\mathbf{w}}(\mathbf{x}) = \sum_{k=0}^{K-1} w_k \phi_k(\mathbf{x}) = \mathbf{w}^T \phi(\mathbf{x})$$

where ϕ_k are called basis functions

- The best \mathbf{w} is considered the one which minimizes the sum-squared error over the training data:

$$\sum_{i=1}^m (y_i - h_{\mathbf{w}}(\mathbf{x}_i))^2$$

- We can find the best \mathbf{w} in closed form:

$$\mathbf{w} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{y}$$

or by other methods (e.g. gradient descent - as will be seen later)

Linear models in general

- By linear models, we mean that the hypothesis function $h_{\mathbf{w}}(\mathbf{x})$ is a *linear function of the parameters \mathbf{w}*
- This *does not mean the $h_{\mathbf{w}}(\mathbf{x})$ is a linear function of the input vector \mathbf{x}* (e.g., polynomial regression)
- In general

$$h_{\mathbf{w}}(\mathbf{x}) = \sum_{k=0}^{K-1} w_k \phi_k(\mathbf{x}) = \mathbf{w}^T \phi(\mathbf{x})$$

where ϕ_k are called *basis functions*

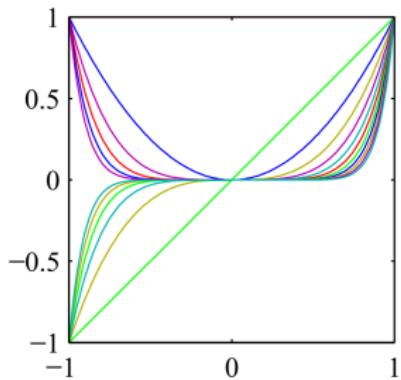
- Usually, we will assume that $\phi_0(\mathbf{x}) = 1, \forall \mathbf{x}$, to create a bias term
- The hypothesis can alternatively be written as:

$$h_{\mathbf{w}}(\mathbf{x}) = \Phi \mathbf{w}$$

where Φ is a matrix with one row per instance; row j contains $\phi(\mathbf{x}_j)$.

- Basis functions are *fixed*

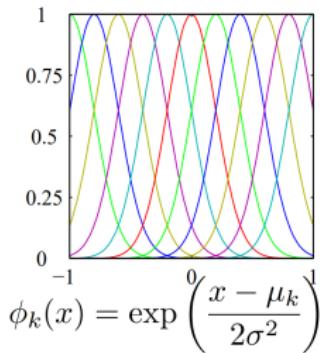
Example basis functions: Polynomial



$$\phi_k(x) = x^k$$

“Global” functions: a small change in x may cause large change in the output of many basis functions

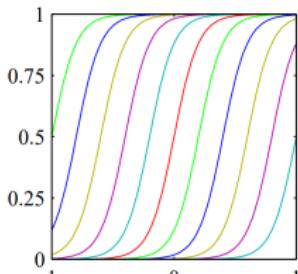
Example basis functions: Gaussian



- μ_k controls the position along the x-axis
- σ controls the width (activation radius)
- μ_k, σ fixed for now (later we discuss adjusting them)
- Usually thought as “local” functions: if σ is relatively small, a small change in x only causes a change in the output of a few basis functions (the ones with means close to x)

Slide credit: Doina Precup

Example basis functions: Sigmoidal



$$\phi_k(x) = \sigma\left(\frac{x - \mu_k}{s}\right) \text{ where } \sigma(a) = \frac{1}{1 + \exp(-a)}$$

- μ_k controls the position along the x-axis
- s controls the slope
- μ_k, s fixed for now (later we discuss adjusting them)
- “Local” functions: a small change in x only causes a change in the output of a few basis (most others will stay close to 0 or 1)

Slide credit: Doina Precup

Polynomial regression

- Given data: $(x_1, y_1), (x_2, y_2), \dots, (x_m, y_m)$.
- Suppose we want a degree- d polynomial fit.
- Let \mathbf{y} be as before and let

$$\mathbf{X} = \begin{bmatrix} x_1^d & \dots & x_1^2 & x_1 & 1 \\ x_2^d & \dots & x_2^2 & x_2 & 1 \\ \vdots & & \vdots & \vdots & \vdots \\ x_m^d & \dots & x_m^2 & x_m & 1 \end{bmatrix}$$

- Solve the linear regression $\mathbf{X}\mathbf{w} \approx \mathbf{y}$.

Example: second-order polynomial regression

$$\mathbf{X} = \begin{bmatrix} 0.75 & 0.86 & 1 \\ 0.01 & 0.09 & 1 \\ 0.73 & -0.85 & 1 \\ 0.76 & 0.87 & 1 \\ 0.19 & -0.44 & 1 \\ 0.18 & -0.43 & 1 \\ 1.22 & -1.10 & 1 \\ 0.16 & 0.40 & 1 \\ 0.93 & -0.96 & 1 \\ 0.03 & 0.17 & 1 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} 2.49 \\ 0.83 \\ -0.25 \\ 3.10 \\ 0.87 \\ 0.02 \\ -0.12 \\ 1.81 \\ -0.83 \\ 0.43 \end{bmatrix}$$

Slide credit: Doina Precup

Example: second-order polynomial regression

$$\mathbf{X}^T \mathbf{X} =$$
$$\begin{bmatrix} 0.75 & 0.01 & 0.73 & 0.76 & 0.19 & 0.18 & 1.22 & 0.16 & 0.93 & 0.03 \\ 0.86 & 0.09 & -0.85 & 0.87 & -0.44 & -0.43 & -1.10 & 0.40 & -0.96 & 0.17 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \end{bmatrix} \times \begin{bmatrix} 0.75 & 0.86 & 1 \\ 0.01 & 0.09 & 1 \\ 0.73 & -0.85 & 1 \\ 0.76 & 0.87 & 1 \\ 0.19 & -0.44 & 1 \\ 0.18 & -0.43 & 1 \\ 1.22 & -1.10 & 1 \\ 0.16 & 0.40 & 1 \\ 0.93 & -0.96 & 1 \\ 0.03 & 0.17 & 1 \end{bmatrix}$$
$$= \begin{bmatrix} 4.11 & -1.64 & 4.95 \\ -1.64 & 4.95 & -1.39 \\ 4.95 & -1.39 & 10 \end{bmatrix}$$

Slide credit: Doina Precup

Example: second-order polynomial regression

$$\mathbf{X}^T \mathbf{y} =$$
$$\begin{bmatrix} 0.75 & 0.01 & 0.73 & 0.76 & 0.19 & 0.18 & 1.22 & 0.16 & 0.93 & 0.03 \\ 0.86 & 0.09 & -0.85 & 0.87 & -0.44 & -0.43 & -1.10 & 0.40 & -0.96 & 0.17 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \end{bmatrix} \times \begin{bmatrix} 2.49 \\ 0.83 \\ -0.25 \\ 3.10 \\ 0.87 \\ 0.02 \\ -0.12 \\ 1.81 \\ -0.83 \\ 0.43 \end{bmatrix}$$
$$= \begin{bmatrix} 3.60 \\ 6.49 \\ 8.34 \end{bmatrix}$$

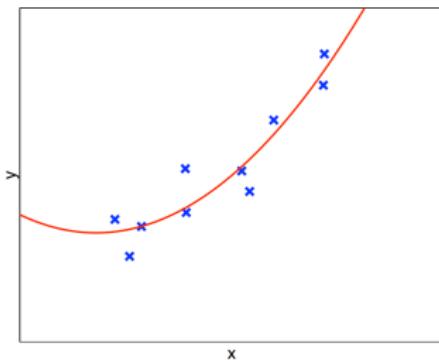
Slide credit: Doina Precup

Example: second-order polynomial regression

$$\begin{aligned}\mathbf{w} &= (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \\ &= \begin{bmatrix} 4.11 & -1.64 & 4.95 \\ -1.64 & 4.95 & -1.39 \\ 4.95 & -1.39 & 10 \end{bmatrix}^{-1} \begin{bmatrix} 3.60 \\ 6.49 \\ 8.34 \end{bmatrix} = \begin{bmatrix} 0.68 \\ 1.74 \\ 0.73 \end{bmatrix}\end{aligned}$$

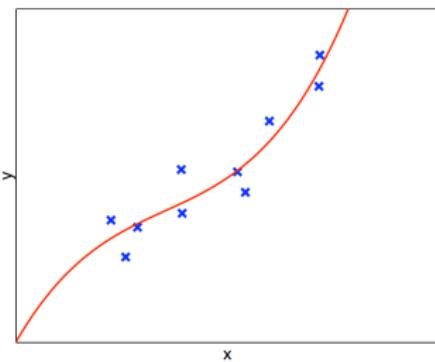
So the best order-2 polynomial is $y = 0.68x^2 + 1.74x + 0.73$.

Higher-order polynomial regression: Order-2 fit



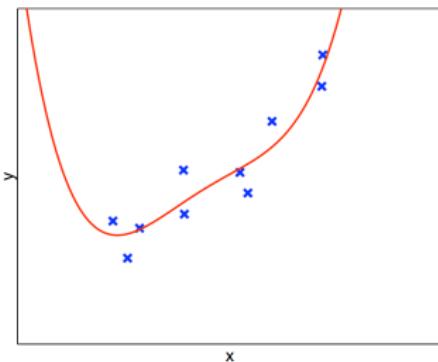
Is this a better fit to the data?

Higher-order polynomial regression: Order-3 fit



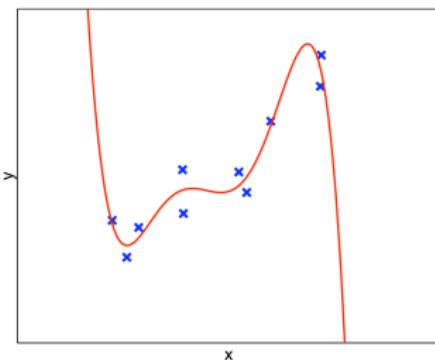
Is this a better fit to the data?

Higher-order polynomial regression: Order-4 fit



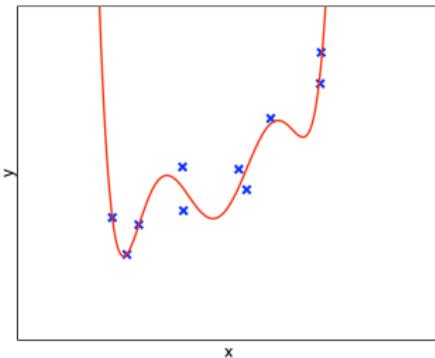
Is this a better fit to the data?

Higher-order polynomial regression: Order-5 fit



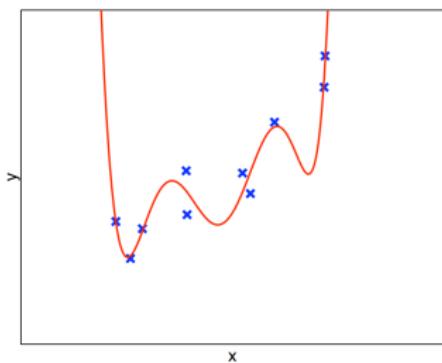
Is this a better fit to the data?

Higher-order polynomial regression: Order-6 fit



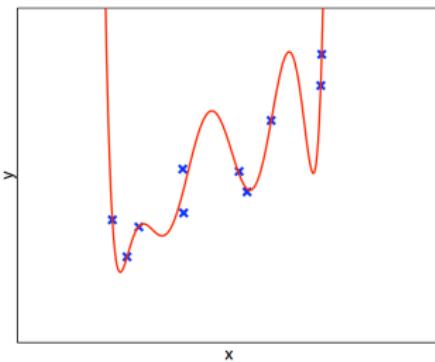
Is this a better fit to the data?

Higher-order polynomial regression: Order-7 fit



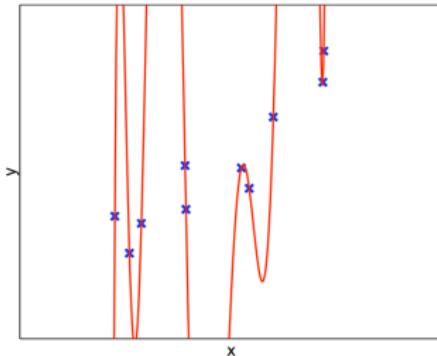
Is this a better fit to the data?

Higher-order polynomial regression: Order-8 fit



Is this a better fit to the data?

Higher-order polynomial regression: Order-9 fit



Is this a better fit to the data?

Overfitting and model selection

- ▶ As long as the hypothesis is complicated enough, we can always achieve zero training loss (e.g., memorize all the data)
- ▶ Minimizing the training loss does NOT indicate a good test performance
- ▶ **Overfitting:** A (complicated) model fits data too well, but cannot generalize at all to new data
- ▶ An extremely important problem for (almost all) machine learning algorithms