

Artificial Intelligence II (CS4442B & CS9542B)

Overfitting, Cross-Validation, and Regularization

Boyu Wang
Department of Computer Science
University of Western Ontario

Motivation examples: polynomial regression

- As the degree of the polynomial increases, there is more degrees of freedom, and the (training) error approaches to zero.

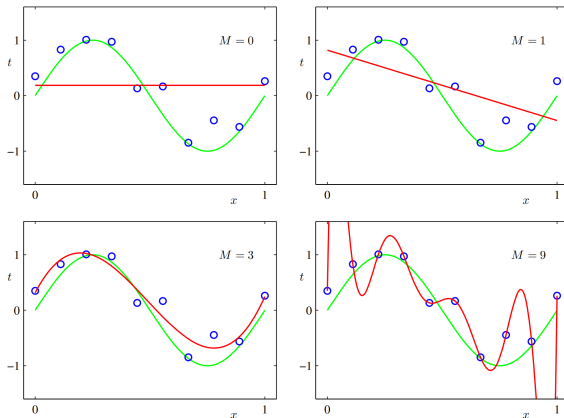



Figure credit: Christopher Bishop

Motivation examples: polynomial regression

- ▶ Minimizing the **training/empirical** loss does NOT indicate a good **test/generalization** performance.
- ▶ **Overfitting**: Very low training error, very high test error.



	$M = 0$	$M = 1$	$M = 6$	$M = 9$
w_0^*	0.19	0.82	0.31	0.35
w_1^*		-1.27	7.99	232.37
w_2^*			-25.43	-5321.83
w_3^*			17.37	48568.31
w_4^*				-231639.30
w_5^*				640042.26
w_6^*				-1061800.52
w_7^*				1042400.18
w_8^*				-557682.99
w_9^*				125201.43

Figure credit: Christopher Bishop

Overfitting – general phenomenon

- ▶ Too simple (e.g., small M) \rightarrow underfitting
- ▶ Too complex (e.g., large M) \rightarrow overfitting

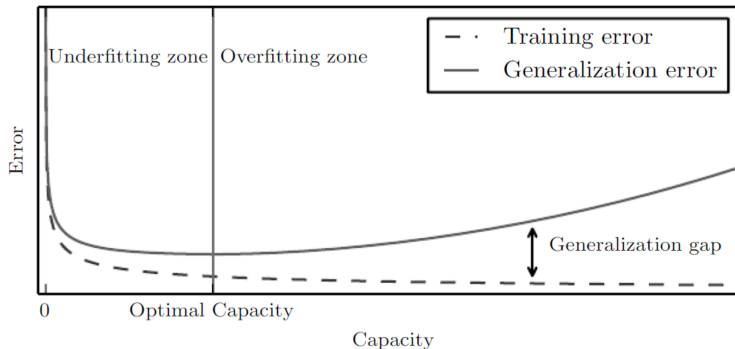


Figure credit: Ian Goodfellow

Overfitting

- ▶ Training loss and test loss are different
- ▶ Larger the hypothesis class, easier to find a hypothesis that fits the training data
 - but may have large test error (overfitting)
- ▶ Prevent overfitting:
 - Large data set
 - Throw away useless hypothesis class (model selection)
 - Control model complexity (regularization)

Larger data set

- ▶ Overfitting is mostly due to sparseness of data.
- ▶ Same model complexity: more data \Rightarrow less overfitting. With more data, more complex (i.e. more flexible) models can be used.

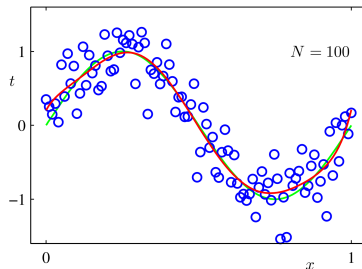
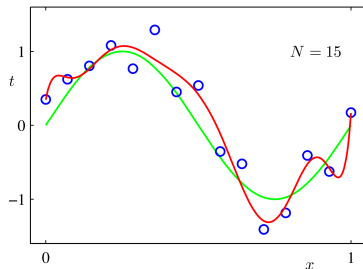


Figure credit: Christopher Bishop

Model selection

- ▶ How to choose the optimal model complexity/**hyper-parameter** (e.g., choose the best degree for polynomial regression)
- ▶ Cannot be done by training data alone

Model selection

- ▶ How to choose the optimal model complexity/**hyper-parameter** (e.g., choose the best degree for polynomial regression)
- ▶ Cannot be done by training data alone
- ▶ We can use our prior knowledge or expertise (e.g., somehow we know that the degree should not exceed 4)
- ▶ Create held-out data to approximate the test error (i.e., mimic the test data)
 - ▶ called **validation data set**

Model selection: cross-validation

For each order of polynomial M

1. Randomly split the training data into K groups, and following procedure K times:
 - i. Leave out the k -th group from the training set as a validation set
 - ii. Use the other other $K - 1$ to find best parameter vector w_k
 - iii. Measure the error of w_k on the validation set; call this J_k
2. Compute the average errors: $J = \frac{1}{K} \sum_{k=1}^K J_k$

Choose the order of polynomial M with the lowest error J .

Model selection: cross-validation




Figure: K -fold cross-validation for the case of $K = 4$

Figure credit: Christopher Bishop

General learning procedure

Given a **training set** and a **test set**

1. Use cross-validation to choose the hyper-parameter/hypothesis class.
2. Once the hyper-parameter is selected, use the entire **training set** to find the best model parameters w .
3. Evaluate the performance of w on the **test set**.

These sets must be disjoint! – you should **never** touch the test data before you evaluate your model.

Summary of cross-validation

- ▶ Can also used for selecting other hyper-parameters for model/algorithm (e.g., number of hidden layers of neural networks, learning rate of gradient descent, or even different machine learning models)
- ▶ Very straightforward to implement algorithm
- ▶ Provides a great estimate of the true error of a model
- ▶ **Leave-one-out cross-validation**: number of groups = number of training instances
- ▶ **Computationally expensive**; even worse when there are more hyper-parameters

Regularization

- ▶ Intuition: complicated hypotheses lead to overfitting
- ▶ Idea: penalize the model complexity (e.g., large values of w_j):

$$L(w) = J(w) + \lambda R(w)$$

where $J(w)$: training loss, $R(w)$: regularization function/regularizer, and $\lambda \geq 0$: regularization parameter to control the tradeoff between data fitting and model complexity.

ℓ_2 -norm regularization for linear regression

Objective function:

$$L(w) = \frac{1}{2} \sum_{i=1}^m \left(\sum_{j=1}^n w_0 + w_j \cdot x_{i,j} - y_i \right)^2 + \frac{\lambda}{2} \sum_{j=1}^n w_j^2$$

► No regularization on w_0 !

Equivalently, we have

$$L(w) = \frac{1}{2} \|Xw - y\|_2^2 + \frac{\lambda}{2} w^\top \hat{I} w$$

where $w = [w_0, w_1, \dots, w_n]^\top$

$$\hat{I} = \begin{bmatrix} 0 & 0 & \dots & 0 \\ 0 & 1 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & 1 \end{bmatrix}$$

ℓ_2 -norm regularization for linear regression

Objective function:

$$\begin{aligned} L(w) &= \frac{1}{2} \|Xw - y\|_2^2 + \frac{\lambda}{2} w^\top \hat{I} w \\ &= \frac{1}{2} \left(w^\top (X^\top X + \lambda \hat{I}) w - w^\top X^\top y - y^\top Xw + y^\top y \right) \end{aligned}$$


Optimal solution (by solving $\nabla L(w) = 0$):

$$w = (X^\top X + \lambda \hat{I})^{-1} X^\top y$$

More on ℓ_2 -norm regularization

$$\arg \min_w \frac{1}{2} \|Xw - y\|_2^2 + \frac{\lambda}{2} w^\top \hat{I} w = (X^\top X + \lambda \hat{I})^{-1} X^\top y$$

- ▶ ℓ_2 -norm regularization pushes the parameters towards to 0.
- ▶ $\lambda = 0 \Rightarrow$ same as in the regular linear regression
- ▶ $\lambda \rightarrow \infty \Rightarrow w \rightarrow 0$
- ▶ $0 < \lambda < \infty \Rightarrow$ magnitude of the weights will be smaller than in the regular linear regression



Another view of ℓ_2 -norm regularization

- From the optimization theory¹, we know that

$$\min_w J(w) + \lambda R(w)$$

is equivalent to

$$\begin{aligned} & \min_w J(w) \\ & \text{such that } R(w) \leq \eta \end{aligned}$$

for some $\eta \geq 0$.

- Hence, ℓ_2 -regularized linear regression can be re-formulated as (we only consider $w_j, j > 0$ here)

$$\begin{aligned} & \min_w \|Xw - y\|_2^2 \\ & \text{such that } \|w\|_2^2 \leq \eta \end{aligned}$$

¹e.g., Boyd and Lieven. Convex Optimization. 2004.

Visualizing ℓ_2 -norm regularization (2 features)

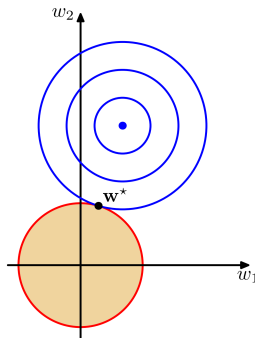


Figure: $w^* = (X^\top X + \lambda I)^{-1} Xy$

Figure credit: Christopher Bishop

ℓ_1 -norm regularization

- Instead of using ℓ_2 -norm, we use ℓ_1 -norm to control the model complexity:

$$\min_w \frac{1}{2} \sum_{i=1}^m \left(\sum_{j=1}^n w_0 + w_j \cdot x_{i,j} - y_i \right)^2 + \lambda \sum_{j=1}^n |w_j|$$

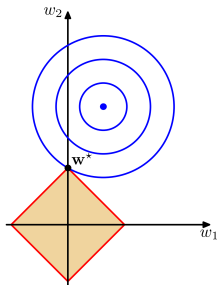
which is equivalent to

$$\begin{aligned} \min_w \quad & \frac{1}{2} \sum_{i=1}^m \left(\sum_{j=1}^n w_0 + w_j \cdot x_{i,j} - y_i \right)^2 \\ \text{such that} \quad & \sum_{j=1}^n |w_j| \leq \eta \end{aligned}$$

- Also called **LASSO** (least absolute shrinkage and selection operator).
- No analytical solution anymore!

Visualizing ℓ_1 -norm regularization (2 features)

- ▶ If λ is large enough, the circle is very likely to intersect the diamond at one of the corners.
- ▶ This makes ℓ_1 -norm regularization much more likely to make some weights exactly 0.
- ▶ In other words, we essentially perform **feature selection**!



Comparison of ℓ_2 and ℓ_1

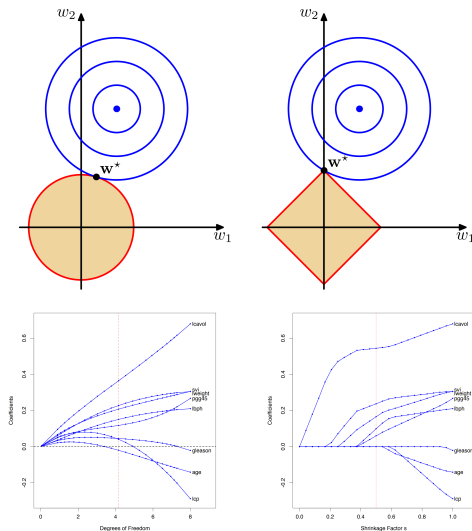


Figure credit: Bishop; Hastie, Tibshirani & Friedman

Summary of regularization

- ▶ Both are commonly used approaches to avoid overfitting.
- ▶ Both push the weights towards 0.
- ▶ ℓ_2 produces small, but non-zero weights, while ℓ_1 is likely to make some weights exactly 0.
- ▶ ℓ_1 optimization is computationally more expensive than ℓ_2 .
- ▶ Choose appropriate λ : cross-validation is often used.