

Heuristic Search

Alice Gao

Lecture 4

Readings: RN 3.5 (esp. 3.5.2), PM 3.6

Outline

Learning Goals

Why Heuristic Search

Greedy Best-First Search

A* Search

Designing an Admissible Heuristic

Pruning the Search Space

Learning goals

By the end of the lecture, you should be able to

- ▶ Describe motivations for applying heuristic search algorithms.
- ▶ Trace the execution of and implement the Greedy best-first search and A* search algorithm with a given heuristic function.
- ▶ Describe properties of the Greedy best-first and A* search.
- ▶ Design an admissible heuristic function for a search problem.
- ▶ Describe reasons for choosing one heuristic function over another.

Learning Goals

Why Heuristic Search

Greedy Best-First Search

A* Search

Designing an Admissible Heuristic

Pruning the Search Space

Why Heuristic Search?

How would ____ choose which one of the two states to expand?

- ▶ an uninformed search algorithm
- ▶ humans

| | | |
|---|---|---|
| 5 | 3 | |
| 8 | 7 | 6 |
| 2 | 4 | 1 |

| | | |
|---|---|---|
| 1 | 2 | 3 |
| 4 | 5 | |
| 7 | 8 | 6 |

Why Heuristic Search

An uninformed search algorithm

- ▶ considers every state to be the same.
- ▶ does not know which state is closer to the goal.
- ▶ may not find the optimal solution.

An heuristic search algorithm *→ using domain knowledge.*

- ▶ uses heuristics to estimate how close the state is to a goal.
- ▶ try to find the optimal solution. *(also find a solution faster.)*

The Heuristic Function

Definition (search heuristic)

*best/
shortest*

A **search heuristic** $h(n)$ is an estimate of the cost of the **cheapest** path from node n to a goal node.

- ▶ $h(n)$ is arbitrary, non-negative, and problem-specific.
- ▶ If n is a goal node, $h(n) = 0$.
- ▶ $h(n)$ must be easy to compute (without search).

Learning Goals

Why Heuristic Search

Greedy Best-First Search

A* Search

Designing an Admissible Heuristic


Pruning the Search Space

Greedy Best-First Search

- ▶ Frontier is a priority queue ordered by the heuristic $h(n)$.
- ▶ Expand the node with the lowest $h(n)$.

Greedy BFS: will it find a solution/terminate?

h : Euclidean straight line distance




$h(x) < h(A)$
for $x = B, C, D, E, F$

Greedy search gets stuck in the loop and never terminates.

Q: What about LCFS?

Greedy BFS: will it find the optimal solution?

frontier: S, SA, SB, \cancel{SAG} solution found: SAG
 $h(A)=6$ $h(B)=8$ $h(G)=0$



What went wrong?

① ignored the costs.

② h values are
deceiving.

optimal solution: $SBCG$

$\text{cost}(SAG) = 21$ $\text{cost}(SBCG) = 11$

Try Greedy Search on the grid search problem

- ▶ Number the nodes as they are removed from the frontier.
- ▶ Use multi-path pruning.
- ▶ Use Manhattan distance as our heuristic function.

tie breaking order : up, left, right, down.



Learning Goals

Why Heuristic Search


Greedy Best-First Search

A* Search

Designing an Admissible Heuristic

Pruning the Search Space

A* Search



- ▶ The frontier is a priority queue ordered by $cost(n) + h(n)$.
- ▶ Expand the node with the lowest $f(n)$.
- ▶ A mix of lowest-cost-first and greedy best-first search.
- ▶ Selects the node on the frontier with the lowest estimated distance from the start to a goal node constrained to go via that node.

Try A* Search on the grid search problem

- ▶ Number the nodes as they are removed from the frontier.
- ▶ Use multi-path pruning.
- ▶ Use Manhattan distance as our heuristic function.

tie breaking rule: up, left, right, down. remove most recently added node first.

| | | | | | | | |
|------------|------------|------------|------------|------------|------------|------------|-------------|
| 14 4 10 | 14 3 11 | 14 2 12 | 16 3 13 | 16 4 12 | 16 5 11 | 16 6 10 | 16 7 9 |
| 12 3 9 | 12 2 10 | 12 1 11 | 14 2 12 | 14 3 11 | 14 4 10 | 14 5 9 | 14 6 8 |
| 10 2 8 | 10 1 22 | 10 0 9 | 12 1 23 | 12 2 10 | 12 3 9 | 12 4 8 | 12 5 7 |
| 10 3 7 | 10 2 21 | 10 1 8 | | | | | 12 6 6 |
| 10 4 6 | | | 4 2 2 | 4 3 3 | 6 4 2 | 8 5 3 | 10 6 14 |
| 10 5 5 | 8 4 19 | 4 3 4 | | | 4 4 0 | 6 5 1 | 8 6 5 |
| 10 6 4 | 8 5 3 | 6 4 8 | 6 5 2 | 8 6 7 | 10 7 1 | 12 8 6 | 14 9 5 |
| 12 7 5 | 10 6 4 | 8 5 18 | 6 4 3 | 10 7 2 | 12 8 11 | 14 9 3 | 16 10 17 |

f heuristic
order
of
cost expansion

A* is Optimal

If the heuristic $h(n)$ is admissible,
the solution found by A* is optimal.

Definition (admissible heuristic)

A heuristic $h(n)$ is **admissible** if it is never an overestimate of the cost of the cheapest path from node n to a goal node.

$h^*(n)$ = the actual cost of
the cheapest path
from n to a goal node.

0 admissible.

$$0 \leq h(n) \leq h^*(n)$$

A* is Optimally Efficient

Among all optimal algorithms that start from the same start node and use the same heuristic, A* expands the fewest nodes.

Learning Goals

Why Heuristic Search

Greedy Best-First Search

A* Search

Designing an Admissible Heuristic

Pruning the Search Space

Some Heuristic Functions for 8-Puzzle

- ▶ Manhattan Distance Heuristic: h_1
The sum of the Manhattan distances of the tiles from their goal positions
- ▶ Misplaced Tile Heuristic: h_2
The number of tiles that are NOT in their goal positions

Both heuristic functions are admissible.

Initial State (S) $h_2(S) = 7$

| | | |
|----------------|----------------|----------------|
| 5 ¹ | 3 ² | ³ |
| 8 ⁴ | 7 ⁵ | 6 ⁶ |
| 2 ⁷ | 4 ⁸ | 1 |

$h_1(S)$
 $= 4 + 3$
 $+ 1 + 2$
 $+ 2 + 0$
 $+ 2 + 2$
 $= 16$

Goal State

| | | |
|---|---|---|
| 1 | 2 | 3 |
| 4 | 5 | 6 |
| 7 | 8 | |

Constructing an Admissible Heuristic

- ▶ Define a relaxed problem by simplifying or removing constraints on the original problem.
- ▶ Solve the relaxed problem without search.
- ▶ The cost of the optimal solution to the relaxed problem is an admissible heuristic for the original problem.

Constructing an Admissible Heuristic for 8-Puzzle

8-puzzle: A tile can move from square A to square B

- ▶ if A and B are adjacent, and
- ▶ B is empty.

Which heuristics can we derive from relaxed versions of this problem?

CQ: Constructing an Admissible Heuristic

CQ: Which heuristics can we derive from the following relaxed 8-puzzle problem?


A tile can move from square A to square B if A and B are adjacent. *(B may not be empty.)*

(A) The Manhattan distance heuristic

(B) The Misplaced tile heuristic

(C) Another heuristic not described above

$$4 + 3 + 1 + 2 + \dots$$



CQ: Constructing an Admissible Heuristic

CQ: Which heuristics can we derive from the following relaxed 8-puzzle problem? *(A and B may not be adjacent.)*


A tile can move from square A to square B. *(B may not be empty.)*

(A) The Manhattan distance heuristic


(B) The Misplaced tile heuristic

(C) Another heuristic not described above

$$\begin{aligned} &1+1+1+1 \\ &+ \dots + \end{aligned}$$



Which Heuristic is Better?



- ▶ We want a heuristic to be admissible.
- ▶ Prefer a heuristic that is very different for different states.
- ▶ Want a heuristic to have higher values (close to h^*).

Dominating Heuristic

Definition (dominating heuristic)

Given heuristics $h_1(n)$ and $h_2(n)$. $h_2(n)$ dominates $h_1(n)$ if

- ▶ $(\forall n (h_2(n) \geq h_1(n)))$.
- ▶ $(\exists n (h_2(n) > h_1(n)))$.


Theorem

If $h_2(n)$ dominates $h_1(n)$, A^ using h_2 will never expand more nodes than A^* using h_1 .*

CQ: Which Heuristic of 8-puzzle is Better?

CQ: Which of the two heuristics of the 8-puzzle is better?

- (A) The Manhattan distance heuristic dominates the Misplaced tile heuristic.
- (B) The Misplaced tile heuristic dominates the Manhattan distance heuristic.
- (C) I don't know....



Learning Goals

Why Heuristic Search

Greedy Best-First Search


A* Search

Designing an Admissible Heuristic

Pruning the Search Space


Cycle Pruning

- ▶ A cycle cannot be part of a least-cost path.
- ▶ Works well with depth-first search.
- ▶ The complexity of cycle pruning is ...



Multiple-Path Pruning

- ▶ If we have already found a path to a node, we can prune other paths to the same node.
- ▶ Subsumes a cycle check.
- ▶ Requires storing all nodes we have found paths to.
- ▶ What if a subsequent path to n is shorter than the first path found?




Lowest-cost-first search w/ multiple-path pruning

If we perform multiple-path pruning with lowest-cost-first search, is it possible for us to prune the optimal solution (least-cost path)?

(A) Yes, it is possible.

(B) No, it is not possible.

$$\text{cost}(P_1) > \text{cost}(P_2)$$



A* search w/ multiple-path pruning


Q: come up w/ a search graph s.t.

A* w/ multi-path pruning discards
the optimal solution.

If we perform multiple-path pruning with A* search,
is it possible for us to prune the optimal solution (least-cost path)?

(A) Yes, it is possible.

(B) No, it is not possible.



Find Optimal Solution w/ Multiple-Path Pruning

What if a subsequent path to n is shorter than the first path found?

- ▶ Remove all paths from the frontier that use the longer path.
- ▶ Change the initial segment of the paths on the frontier to use the shorter path.
- ▶ Make sure that the least-cost path to a node is found first.

A* search w/ multiple-path pruning

How can we ensure that A* with multiple-path pruning is optimal?

- ▶ Ensure that we find the least-cost path to every node first.
- ▶ Admissible heuristic guarantees the above for a goal node, but not for other nodes. *→ consistent.*
- ▶ We need the heuristic to satisfy the monotone restriction:

for any arc $\langle m, n \rangle$, $\underline{h(m)} - \underline{h(n)} \leq \underline{\text{cost}(m, n)}$.

if n is a goal node, $h(m) \leq \text{cost}(m, n) \rightarrow$ admissible heuristic.

If the heuristic satisfies the monotone restriction,

A* search with multiple-path pruning is optimal.

$h(m) - h(n) =$ the heuristic estimate of the path cost from m to n .

Summary of Search Strategies

| Strategy | Frontier Selection | Halts? | Space | Time |
|-------------------|------------------------------|--------|--------|------|
| Depth-first | Last node added | No | Linear | Exp |
| Breadth-first | First node added | Yes | Exp | Exp |
| Lowest-cost-first | $\min \text{cost}(n)$ | Yes | Exp | Exp |
| Greedy Best-first | $\min h(n)$ | No | Exp | Exp |
| A* | $\min \text{cost}(n) + h(n)$ | Yes | Exp | Exp |

Revisiting the learning goals

By the end of the lecture, you should be able to

- ▶ Describe motivations for applying heuristic search algorithms.
- ▶ Trace the execution of and implement the Greedy best-first search and A* search algorithm with a given heuristic function.
- ▶ Describe properties of the Greedy best-first and A* search.
- ▶ Design an admissible heuristic function for a search problem.
- ▶ Describe strategies for choosing among multiple heuristic functions.
- ▶ Describes strategies for pruning a search space.