

# Local Search

Alice Gao

Lecture 6

Readings: RN 4.1, PM 4.7 - 4.8

# Outline

Learning Goals

Introduction to Local Search

Local Search Algorithms

- Iterative Best Improvement

- Greedy descent with randomization

- Simulated Annealing

- Population-Based Algorithms

Revisiting the Learning goals

# Learning Goals

By the end of the lecture, you should be able to

- ▶ Describe the advantages of local search over other search algorithms.
- ▶ Formulate a real world problem as a local search problem.
- ▶ Verify whether a state is a local/global optimum.
- ▶ Describe strategies for escaping local optima.
- ▶ Trace the execution of greedy descent, greedy descent with random restarts, simulated annealing, and genetic algorithms.
- ▶ Compare and contrast the properties of local search algorithms.

# Why use local search?

So far, the search algorithms explore the space systematically and keep track of one or more paths.

- ▶ The search space can be too big or infinite.
- ▶ For solving CSPs, the path to a goal is irrelevant.

Solution: local search

# Properties of local search

- ▶ Explores only a portion of the search space.
- ▶ Requires little memory.

## Advantages:

- ▶ Can find solutions quickly on average.
- ▶ Works for CSPs and general optimization problems.

## Disadvantages:

- ▶ No guarantee that a solution will be found if one exists.  
Cannot prove that no solution exists.

# What is local search?

- ▶ Start with a complete assignment of values to variables.
- ▶ Take steps to improve the solution iteratively.

# Local Search

A local search problem consists of:

- ▶ A state : a complete assignment to *all* of the variables.
- ▶ A neighbour relation: which states do I explore next?
- ▶ A cost function: how good is each state?

## 4-Queens Problem as a Local Search Problem

Variables:  $x_0, x_1, x_2, x_3$  where  $x_i$  is the row position of the queen in column  $i$ . Assume that there is one queen per column.  $i \in \{0, 1, 2, 3\}$ .

Domains:  $x_i \in \{0, 1, 2, 3\} \forall i$ . row positions.

Initial state: 4 queens on the board in random  $\wedge$

Goal state: 4 queens on the board.

No pair of queens are attacking each other.


Neighbour relation:

- Version A: move a single queen to a different row in the same column.
- Version B: swap the row positions of two queens.



## 4-Queens Problem as a Local Search Problem

Cost function: the number of pairs of queens attacking each other, directly or indirectly.



cost = 4

Learning Goals

Introduction to Local Search

Local Search Algorithms

- Iterative Best Improvement

- Greedy descent with randomization

- Simulated Annealing

- Population-Based Algorithms

Revisiting the Learning goals

## Greedy descent

If a neighbour is an improvement,  
then move to the best neighbour.

Greedy: does not look beyond the immediate neighbours.

a.k.a. hill climbing or greedy ascent.

- ▶ Start with a random state.
- ▶ Move to a neighbour with the lowest cost if it's better than the current state.
- ▶ Stop when no neighbour has a lower cost than current state.
- ▶ Only remembers the current node.

requires <sup>↓</sup> very  
little memory.

If no neighbour is an improvement, stop!  
current state is one of the best  
states in the neighbourhood.

## Greedy descent in one sentence

- ① descend: minimize cost.  
move to a neighbour w/ a lower cost.
- ② thick fog: can only see immediate neighbours.

① Descend into a canyon in a thick fog with amnesia

- ③ amnesia : no memory of where we've been.  
may stumble on the same state multiple  
times w/ realizing it.

# Properties of Greedy Descent

- ▶ Performs quite well in practice.  
Makes rapid progress towards a solution.


*the state w/ lowest cost  
among all the states.*

- ▶ Given enough time,  
will greedy descent find the global optimum?

*No. Greedy descent can get  
stuck at local optimums.*

# Where can Greedy Descent get stuck?

- ▶ Local optimum: No neighbour has a (strictly) lower cost.
- ▶ Global optimum: A state that has the lowest cost among all the states. *(a special type of local optimum.)*



## CQ: Local and global optimum (1)

**CQ:** Consider the following state of the 4-queens problem. Consider neighbour relation B: swap the row positions of two queens. Which of the following is correct?

state : 3201 (2)

2301 (4)

0231 (1)

1203 (1)

0			Q	
1				Q
2		Q		
3	Q			

3021 (1)

3102 (1)

3210 (6)

- (A) This state is a local optimum and is a global optimum.
- (B) This state is a local optimum and is NOT a global optimum.
- (C) This state is NOT a local optimum and NOT a global optimum.

## CQ: Local and global optimum (2)

**CQ:** Consider the following state of the 4-queens problem. Consider neighbour relation A: move a single queen to another square in the same column. Which of the following is correct?

2	2	Q	4
3	3	4	Q
3	Q	3	2
Q	3	3	2

*cost = 2*

- (A) This is a local optimum and is a global optimum.
- (B) This is a local optimum and is NOT a global optimum.
- (C) This is NOT a local optimum and NOT a global optimum.



# Escaping flat local optimums

*stop after a # of consecutive sideways moves*

- ▶ Sideway moves: allow the algorithm to move to a neighbour that has the same cost.
- ▶ Tabu list: keep a small list of recently visited states and forbid the algorithm to return to those states.

*some short term memory.*

# Performance of Greedy Descent with sideway moves

8-queens problem:  $\approx 17$  million states.

- ▶ Greedy descent

% of instances solved: 14%

# of steps until success/failure: 3-4 steps on average until success or failure.

- ▶ Greedy descent +  $\leq 100$  consecutive sideway moves:

% of instances solved: 94%

# of steps until success/failure: 21 steps until success and 64 steps until failure.

Learning Goals

Introduction to Local Search

Local Search Algorithms

- Iterative Best Improvement

- Greedy descent with randomization

- Simulated Annealing

- Population-Based Algorithms

Revisiting the Learning goals

# Random restarts and random walks

Greedy descent can get stuck at a local optimum (that is not a global optimum). What can we do?

- ▶ Random restarts:  
restart search in a different part of the space.  
Example: Greedy descent with random restarts

- ▶ Random walks:  
move to a state with a higher cost occasionally.  
Example: Simulated annealing

# Random restarts vs random walks

(1) Which type of randomization is better for landscape (a) ?


(A) Random restarts *a global random move.*

(B) Random walks

(2) Which type of randomization is better for landscape (b) ?

(A) Random restarts

(B) Random walks *a local random move.*



## Greedy descent with random restarts

*If at first you don't succeed, try, try again.*

- ▶ Performs multiple greedy descents from randomly generated initial states.
- ▶ Will greedy descent with random restarts find the global optimum?

*It will find the global optimum w/ probability approaching 1. Eventually, it will generate the global optimum as the initial state.*

Learning Goals

Introduction to Local Search

Local Search Algorithms

- Iterative Best Improvement

- Greedy descent with randomization

- Simulated Annealing**

- Population-Based Algorithms

Revisiting the Learning goals

So far...

Greedy descent focuses on optimization/exploitation,  
whereas random moves allow us to explore the search space.

Can we combine exploration and optimization into one algorithm?



# Simulated Annealing

- ▶ Annealing: slowly cool down molten metals to make them stronger.
- ▶ Start with a high temperature and reduce it slowly.
- ▶ At each step, choose a random neighbour.
  - If the neighbour is an improvement, move to it.
  - If the neighbour is not an improvement, move to the neighbour probabilistically depending on
    - ▶ the current temperature  $T$
    - ▶ how much worse is the neighbour compared to current state

# How likely do we move to a worse neighbour?

$A$  is the current state and  $A'$  is the worse neighbour.

Let  $\Delta C = \text{cost}(A') - \text{cost}(A)$ . The current temperature is  $T$ .

We move to the neighbour  $A'$  with probability

$$e^{-\frac{\Delta C}{T}}$$

*Gibbs distribution / Boltzmann distribution.*

## CQ: Probability of moving to a worse neighbour

**CQ 1:**  $A$  is the current state and  $A'$  is the worse neighbour.  
Let  $\Delta C = \text{cost}(A') - \text{cost}(A)$ .

As  $T$  decreases, how does the probability of moving to the worse neighbour ( $e^{-\frac{\Delta C}{T}}$ ) change?

(A) As  $T$  decreases, we are more likely to move to the neighbour.

(B) As  $T$  decreases, we are less likely to move to the neighbour.

$$\Delta C > 0, T > 0, \frac{\Delta C}{T} > 0. \quad T \downarrow, \frac{\Delta C}{T} \uparrow, -\frac{\Delta C}{T} \downarrow, e^{-\frac{\Delta C}{T}} \downarrow$$

$$\Delta C = 10 \quad T = 100 \quad e^{-10/100} = e^{-0.1} = 0.9$$

$$T = 10 \quad e^{-10/10} = e^{-1} = 0.36$$

## CQ: Probability of moving to a worse neighbour

**CQ 2:**  $A$  is the current state and  $A'$  is the worse neighbour.  
Let  $\Delta C = \text{cost}(A') - \text{cost}(A)$ .

As  $\Delta C$  increases (the neighbour becomes worse), how does the probability of moving to the worse neighbour ( $e^{-\frac{\Delta C}{T}}$ ) change?

(A) As  $\Delta C$  increases, we are more likely to move to the neighbour.

(B) As  $\Delta C$  increases, we are less likely to move to the neighbour.

$$\Delta C > 0, T > 0, \frac{\Delta C}{T} > 0. \quad \Delta C \uparrow, \frac{\Delta C}{T} \uparrow, -\frac{\Delta C}{T} \downarrow, e^{-\frac{\Delta C}{T}} \downarrow$$

$$T = 10 \quad \Delta C = 10 \quad e^{-10/10} = 0.36.$$

$$\Delta C = 100 \quad e^{-100/10} = 0.000045.$$

# Simulated Annealing Algorithm


---

## Algorithm 1 Simulated Annealing

---

```
1: current  $\leftarrow$  initial-state
2:  $T \leftarrow$  a large positive value
3: while  $T > 0$  do
4:   next  $\leftarrow$  a random neighbour of current
5:    $\Delta C \leftarrow \text{cost}(\text{next}) - \text{cost}(\text{current})$ 
6:   if  $\Delta C < 0$  then
7:     current  $\leftarrow$  next
8:   else
9:     current  $\leftarrow$  next with probability  $p = e^{\frac{-\Delta C}{T}}$ 
10:  decrease  $T$ 
11: return current
```

---



# Annealing Schedule

How should we decrease  $T$ ?

- ▶ In theory, we want to decrease the temperature very slowly.

*If the temperature decreases slowly enough, simulated annealing is guaranteed to find the global optimum with probability approaching 1.*

- ▶ In practice, a popular schedule is geometric cooling.

*start w/  $T=10$ . multiply by 0.99 at each step.*

Simulated annealing is like life...

Learning Goals

Introduction to Local Search

Local Search Algorithms

- Iterative Best Improvement

- Greedy descent with randomization

- Simulated Annealing

- Population-Based Algorithms

Revisiting the Learning goals



# Beam Search

- ▶ Remember  $k$  states.
- ▶ Choose the  $k$  best states out of **all of the neighbors**.
- ▶  $k$  controls space and parallelism.

① What is beam search when  $k = 1$ ? *greedy descent*

② How is beam search different from  $k$  random restarts in parallel?

③ Are there problems with beam search?

② *w/ random restarts, each search is independent of the others.  
for beam search, useful info is passed among parallel searches.*

③ *suffers from a lack of diversity among the  $k$  states.  
can quickly become concentrated in a small region.*

# Stochastic Beam Search

- ▶ Choose the  $k$  states probabilistically.
- ▶ Probability of choosing a neighbour is proportional to its fitness.
- ▶ Maintains diversity in the population of states.
- ▶ Mimics natural selection.

# Genetic Algorithm

*stochastic beam search – asexual reproduction*  
*genetic algorithm – sexual reproduction.*

- ▶ Maintain a population of  $k$  states.
- ▶ Randomly choose two states to reproduce.  
Probability of choosing a state for reproduction is proportional to the fitness of the state.
- ▶ Two parent states crossover to produce a child state.
- ▶ The child state mutates with a small probability.
- ▶ Repeat the steps above to produce a new population.
- ▶ Repeat until the stopping criteria is satisfied.

# Comparing greedy descent and genetic algorithm

- ▶ How do the algorithms explore the state space?

Greedy descent generates neighbours of the state based on the neighbour relation.

Genetic algorithm ... *crosses over parent states randomly to produce a child state and mutates the child state.*

- ▶ How do the algorithms optimize the quality of the population?

Greedy descent moves to the best neighbour.

Genetic algorithm ... *chooses parent states probabilistically based on the fitness of the states.*

# Revisiting the Learning Goals

By the end of the lecture, you should be able to

- ▶ Describe the advantages of local search over other search algorithms.
- ▶ Formulate a real world problem as a local search problem.
- ▶ Verify whether a state is a local/global optimum.
- ▶ Describe strategies for escaping local optima.
- ▶ Trace the execution of greedy descent, greedy descent with random restarts, simulated annealing, and genetic algorithms.
- ▶ Compare and contrast the properties of local search algorithms.