

# REINFORCEMENT LEARNING

---

FUNDAMENTALS  
+  
APPLICATIONS

WEEK 1

---

BANDIT ALGORITHM

# WEEK 1 : CORE CONCEPTS

---

**State:** A representation of the **agent**'s current situation. This can be a simple ID (State 1, State 2, ...), coordinates, or a set of features. It can include compressed representations of the past (i.e., memory).

**Action:** Selected by the **Agent**, according to their **Policy**, that affect the **Environment**.

**Reward:** The special learning signal emitted from the **Environment** in response to **Agent actions**. Used by the **Agent** to learn about which **actions** and **states** are "good". **Rewards** are usually specified by you.

**Timestep:** For this week, **timesteps** represent discrete **action-reward-update** cycles. Our learner takes an **action**, receives some **reward**, and **updates** its **value estimates**. Then we go to the next **timestep**.

**Run:** One **run** of N **timesteps**. We might conduct 1000 **runs** of 2000 **timesteps**.

# WEEK 1 : CORE CONCEPTS

---

**Learning Rate:** Represented by the symbol  $\alpha \in [0,1]$ , learning rate is the weight on the most recent reward.

**Policy:** Represented by the symbol  $\pi$ , policies are how agents choose their actions. An example policy might be “always choose the action that I think gives me the best reward RIGHT NOW.” Examples of policies from this week are epsilon-greedy and Upper-Confidence-Bound.

# BANDIT

ANY NON-LAWFUL

CHALLENGE: 1/8

MEDIUM HUMANOID (ANY RACE)

SPEED: 30'

(25 XP)

AC: 12 HP: 11

STR	DEX	CON	INT	WIS	CHA
11 (+0)	12 (+1)	12 (+1)	10 (+0)	10 (+0)	10 (+0)

SENSES: PASSIVE PERCEPTION 10  
LANGUAGES: ANY ONE LANGUAGE  
(USUALLY COMMON)

## ACTIONS:

### SCIMITAR.

MELEE WEAPON ATTACK: +3 TO HIT, REACH 5 FT., ONE TARGET. HIT: 4 (1D6 + 1) SLASHING DAMAGE.

### LIGHT CROSSBOW.


RANGED WEAPON ATTACK: +3 TO HIT, RANGE 80 FT./320FT., ONE TARGET. HIT: 5 (1D8 + 1) PIERCING DAMAGE.








# The Multi-Armed Bandit Problem










Reward  
distribution



# EQUATION : ACTION-VALUE

---

The estimated **value** of **action a** at **timestep t**.

$$q_*(a) = \mathbb{E}[R_t | A_t = a]$$

“The **value** of arbitrary **action a**, denoted by  **$q_*(a)$** , is the **expected reward r** given that **a** was selected.”

# EQUATION : ACTION-VALUE

---

The estimated **value** of **action a** at **timestep t**.

$$Q_t(a) = \frac{\text{sum of rewards when } a \text{ taken prior to current timestep } t}{\text{number of times } a \text{ taken prior to timestep } t}$$

# EQUATION : ACTION-VALUE

---

The estimated **value** of **action a** at **timestep t**.

$$Q_t(a) = \frac{1}{N(a)} \sum_{i=1}^{n-1} R_i$$

$Q_t(a)$  : The estimated **value** of **action a** at **timestep t**.

$N(a)$  : The number of times action a was taken prior to **timestep t**.

$R_i$  : The reward received at prior **timestep i**.



# EQUATION: ACTION-VALUE

---

INCREMENTAL UPDATE, SAMPLE AVERAGE

$$Q_{n+1} = Q_n + \frac{1}{n} [R_n - Q_n]$$

# ALGORITHM : SIMPLE BANDIT

---

Initialization:  $Q(a) \leftarrow 0, N(a) \leftarrow 0, k \in \mathbb{Z}^{0+}, N \in \mathbb{Z}^{0+}$

---

For  $t = 1, 2, \dots, n$  DO:

$A \leftarrow \begin{cases} \operatorname{argmax}_a Q(a) & \text{with probability } 1 - \epsilon \\ \text{Random action} & \text{with probability } \epsilon \end{cases}$

Take action  $a$ , receive reward  $R$

$N(a) = N(a) + 1$

$Q(a) = Q(a) + \frac{1}{N(a)} [R - Q(a)]$

# EQUATION: ACTION-VALUE

---

EXPONENTIAL RECENCY-WEIGHTED AVERAGE

$$Q_{n+1} = \alpha [ R_n - Q_n ]$$

$\alpha$  : Learning rate.

# EQUATION: ACTION-VALUE

## EXPONENTIAL RECENCY-WEIGHTED AVERAGE




A exponential recency-weighted average (blue line) applied to price action of Apple stock. Note that this is a *lagging indicator*: it only responds to shifts after some delay (lag). The lower  $\alpha$  is, the more it will lag. Sometimes we want a larger lag so our learner is not overly reactive to extreme draws or short-lived shifts in distribution. Sometimes we want a shorter lag so our learner can react quickly to sudden shifts in distribution.

# EQUATION : UCB

UPPER CONFIDENCE BOUND  $\pi$

$$A_t = \operatorname{argmax} Q_t(a) + c[ \sqrt{\log(t) / N_t(a)} ]$$





# GRADIENT BANDIT

## PREFERENCE UPDATE

$$H_{t+1}(A_t) = H_t(A_t) + \alpha(R_t - \bar{R}_t)(1 - \pi_t(A_t))$$

$$H_{t+1}(a) = H_t(a) - \alpha(R_t - \bar{R}_t)\pi_t(a), \quad \forall a \neq A_t$$

$H_t(a)$  preference value for action  $a$

$\alpha$ : Learning rate

$R_t$ : Reward received at time  $t$ .

$\bar{R}_t$ : Average reward received up to and including time  $t$ .

$\pi_t(a_t)$ : The probability [0,1] of choosing action  $a$  at time  $t$ , as obtained by passing our preference values  $H$  through a softmax function.

# GRADIENT BANDIT

## PREFERENCE UPDATE

$$H_{t+1}(A_t) = H_t(A_t) + \alpha(R_t - \bar{R}_t)(1 - \pi_t(A_t))$$

$$H_{t+1}(a) = H_t(a) - \alpha(R_t - \bar{R}_t)\pi_t(a), \quad \forall a \neq A_t$$

$H_t(a)$  preference value for action  $a$

$\alpha$ : Learning rate

$R_t$ : Reward received at time  $t$ .

$\bar{R}_t$ : Average reward received up to and including time  $t$ .

$\pi_t(a_t)$ : The probability [0,1] of choosing action  $a$  at time  $t$ , as obtained by passing our preference values  $H$  through a softmax function.

# DESIGN THINKING

## K-ARMED BANDIT IMPLEMENTATION

### AGENT

Q: action-value estimates

N: Number of times  
actions chosen in current  
run

### POLICY

Contains method for  
choosing actions

Contains relevant  
hyperparameters  
(e.g., epsilon).

### ENVIRONMENT

K : The number of  
actions.

$Q_*$  : The true expected  
reward of each arm.

Method for accepting  
action choice and  
emitting the reward

Method for resetting  
reward.

### TEST RUNNER

Number of timesteps

Performance histories for  
reward and optimal  
action choices (yes/no).

Visualization

Must be able to hold  
(and compare) multiple  
agents.

Holds ONE environment.

# DESIGN THINKING

---

## K-ARMED BANDIT IMPLEMENTATION

An Agent is assigned to ONE environment at a time (it can interact with one bandit problem at once).

- Agents can learn one problem, then “reset” and learn another.

An Agent has ONE policy.

- Just like Environments Agents can swap out one policy for another.

An Environment is a standalone thing that can be interacted with. It has no notion of Agents or Policies.

A Test Runner can have ONE environment at a time. This is for simplicity; we could have multiple but we have no need of that. Environments can be swapped out.

A Test Runner can have MANY Agents at once. This is so we can compare and contrast the performance of different Agents within a single problem space (Environment).

## TEST RUNNER

Epsilon  
Greedy

$$\epsilon = 0.1$$

Sample  
Average  
Agent

$$\alpha = 0.5$$

Epsilon  
Greedy

$$\epsilon = 0.04$$

Sample  
Average  
Agent

$$\alpha = 0.05$$

Environment

- $k = 10$
- Stationary
- Reward distributions: Gaussian

Sample  
Average  
Agent  
 $\alpha = 1/n$

UCB

$$c = 2$$

Gradient  
Agent

$$\alpha = 0.001$$

Softmax